

COT 6405
ANLYSIS OF ALGORITHMS

Dynamic Programming

Computer & Electrical Engineering and Computer Science Dept.
Florida Atlantic University

Spring 2017

Outline

- DP – introduction
- Weighted Interval Scheduling (KT – chapter 6.1)
- Principles of DP (KT – chapter 6.2)
- Change-making problem (BB – chapter 8.2)
- 0-1 knapsack problem (BB – chapter 8.4)
- Sequence alignment problem (KT – chapter 6.6 & 6.7)

KT book - *Algorithm Design* by J. Kleinberg and Eva Tardos

BB book - *Fundamentals of Algorithms* by Gilles Brassard and Paul Bratley

Dynamic Programming (DP)

- is a technique, not a specific algorithm (like divide-and-conquer)
- applied to optimization problems (maximization or minimization)
- applicable when subproblems are not independent, that is subproblems share subsubproblems. Then DP solves each subproblem only once, stores the result in a table, and reuses it later

Weighted Interval Scheduling

Problem definition:

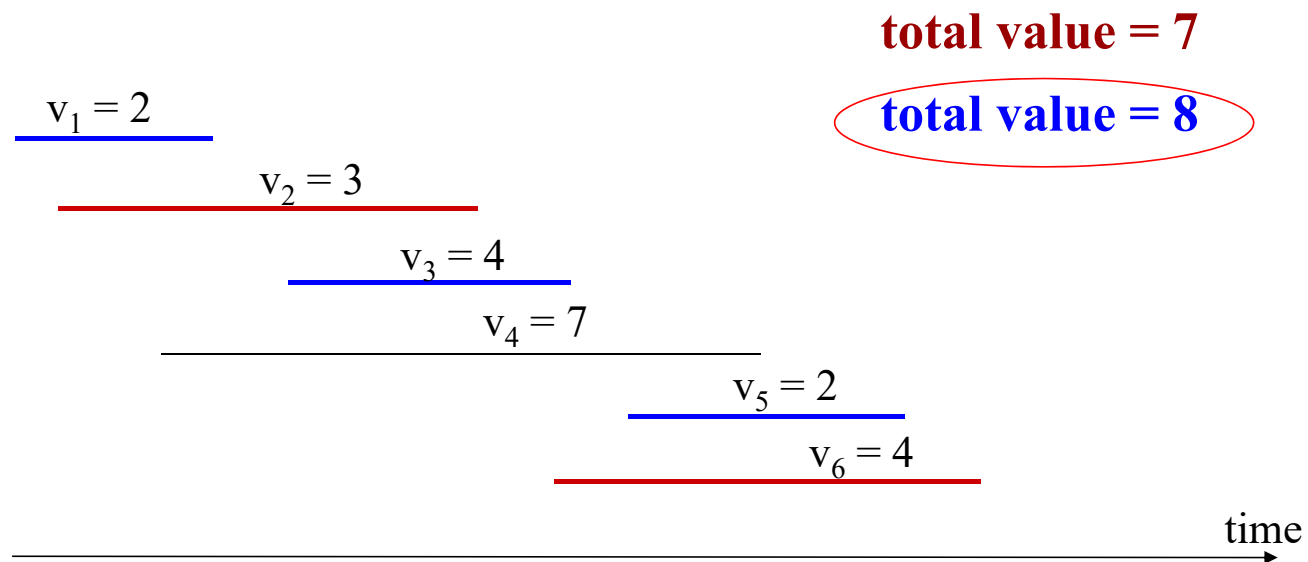
We are given **1 resource** and a number of **n requests** labeled $1, 2, \dots, n$. Each request i has a starting time s_i , a finishing time f_i , and a *value* (or *weight*) $v_i > 0$.

Find a *compatible* subset S of requests (intervals) of *maximum total value* $\sum_{i \in S} v_i$.

- two requests i and j are compatible if the requests do not overlap
 - request i is earlier than j , $f_i \leq s_j$
 - request j is earlier than i , $f_j \leq s_i$
- a subset of requests is compatible if all pairs i, j ($i \neq j$) are compatible

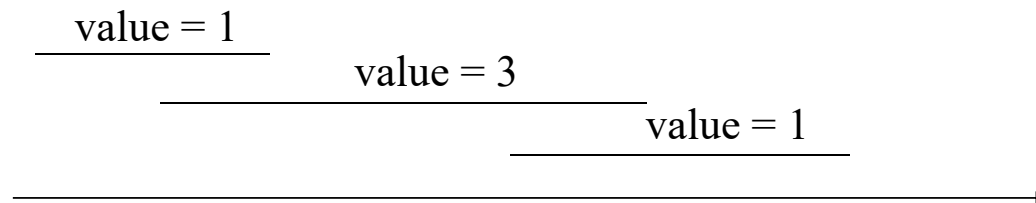
Weighted Interval Scheduling

Example



Weighted Interval Scheduling

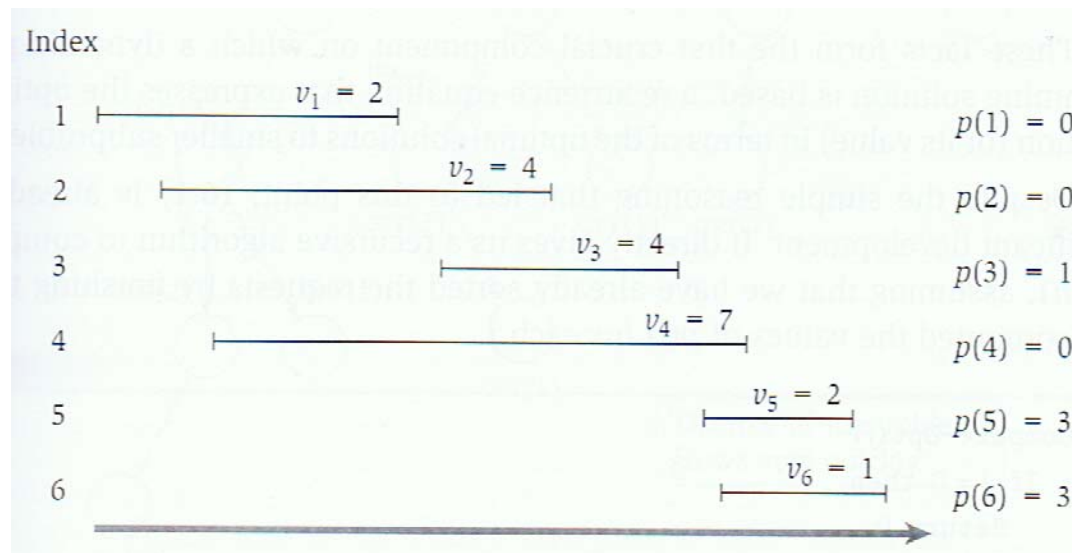
- can be solved using *greedy* when all requests have the same value (equal to 1)
 - greedy choice: choose the compatible request with the earliest finishing time
- when the requests have different values, greedy does not work



- for different request values, DP produces an optimal solution

A recursive algorithm

- suppose that the requests are sorted in nondecreasing finishing time $f_1 \leq f_2 \leq \dots \leq f_n$
- we say that request i *comes before* j if $i < j$
- define $p(j)$ of an interval j – the largest index $i < j$ s.t. i and j are disjoint
 - rightmost interval i that ends before j begins



Weighted Interval Scheduling

- optimal solution O : either n belongs to O or it doesn't
 - if $n \in O$, then in addition O contains an *optimal* solution to the problem with requests $\{1, \dots, p(n)\}$
 - if $n \notin O$, then O is an *optimal* solution to the problem with requests $\{1, 2, \dots, n-1\}$
- let O_j – optimal solution for the problem w/ requests $\{1, 2, \dots, j\}$
 - $OPT(j)$ – value of this solution
- objective: find O_n and $OPT(n)$

Writing a recursion

$$\text{OPT}(j) = \max (v_j + \text{OPT}(p(j)), \text{OPT}(j-1))$$

- request j belongs to an optimal solution on the set $\{1, 2, \dots, j\}$ iff $v_j + \text{OPT}(p(j)) \geq \text{OPT}(j-1)$
- characteristic of DP: write a recurrence equation that expresses the optimal solution (or its value) in terms of optimal solutions to smaller subproblems

Recursive algorithm

Compute-Opt(j)

if $j == 0$ then

 return 0

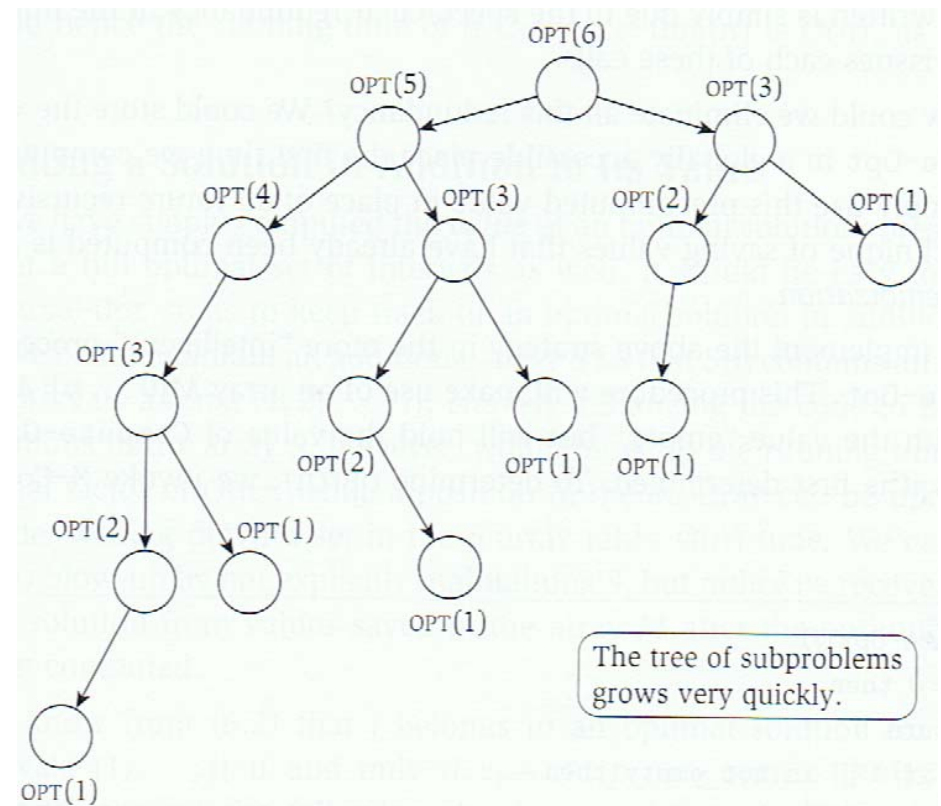
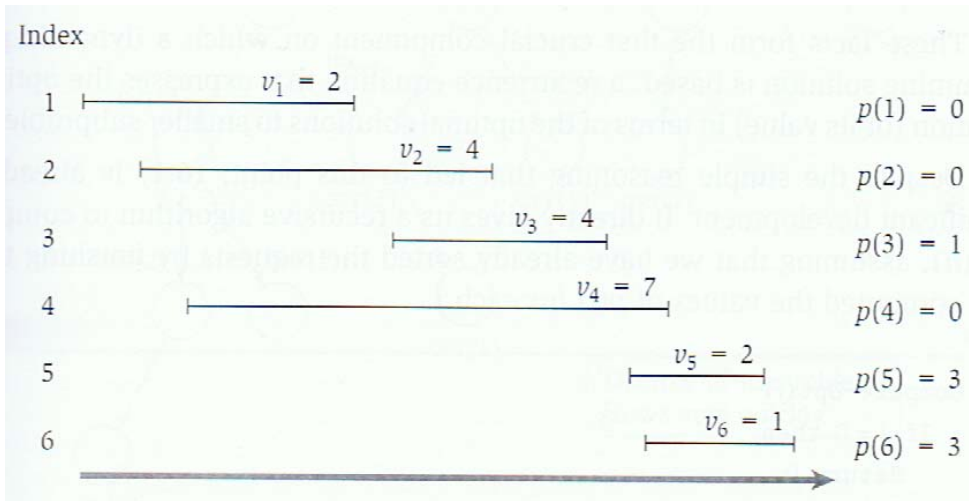
else

 return $\max(v_j + \text{Compute-Opt}(p(j)), \text{Compute-Opt}(j-1))$

endif

- drawback: RT is exponential in n

Example



Memoizing the Recursion

- Key observations:
 - Compute-Opt(n) is only solving $n+1$ different subproblems: Compute-Opt(0), Compute-Opt(1), ..., Compute-Opt(n)
 - exponential time is due to the redundancy in the number of times the same call is made
- **Memoization technique:** *solve each subproblem only once and store its value in a table. All future calls use the precomputed value.*

Memoizing the Recursion

- array $M[0..n]$
 - $M[j]$ initially empty, then store $\text{Compute-Opt}(j)$ value
- to determine $\text{OPT}(n)$, call $\text{M-Compute-Opt}(n)$

M-Compute-Opt(j)

if $j == 0$ then

 return 0

else if $M[j]$ is not empty then

 return $M[j]$

else

$M[j] = \max(v_j + \text{M-Compute-Opt}(p(j)), \text{M-Compute-Opt}(j-1))$

 return $M[j]$

endif

Analyzing the Memoized Version

- Assuming that the input intervals are sorted by their finish time and that the $p()$ values are already computed, the RT to compute $M\text{-Compute-Opt}(n)$ is $O(n)$.
 - Proof: since M has $n+1$ entries, there are at most $O(n)$ calls to $M\text{-Compute-Opt}$, and hence the RT is $O(n)$

A second solution: iterative approach

Iterative-Compute-Opt

$M[0] = 0$

for $j = 1, 2, \dots, n$

$M[j] = \max(v_j + M[p(j)], M[j-1])$

endfor

$RT = O(n)$

- Example

Computing a solution in addition to its value

- Find the subset of requests of maximum value

Find-Solution(j)

```
if j == 0 then
    return  $\emptyset$ 
else
    if  $v_j + M[p(j)] \geq M[j-1]$  then
        output Find-Solution(p(j))  $\cup$  j
    else
        output the result of Find-Solution(j-1)
    endif
endif
```

- initial call: Find-Solution(n)
- RT = $O(n)$, calls recursively on smaller instances and takes constant time per call

Principles of Dynamic Programming

- write solution to a problem recursively, based on solutions to subproblems
- memoization or iteration over subproblems?
 - iteratively building up subproblems is simpler
- properties of DP:
 - there is only a polynomial number of subproblems
 - the solution to the original problem can be easily computed from the solutions to the subproblems
 - there is a natural ordering of subproblems from smallest to largest

Change-making problem

- Unfortunately, even though greedy algorithm is very efficient, it works only in a limited number of instances
- Dynamic programming works for all systems of coins
- Suppose that the currency has available coins of n different denominations
 - a coin of denomination i , $1 \leq i \leq n$ has value $d_i > 0$ units
- Suppose that we have an unlimited supply of coins of each denomination
- Goal: give the customer coins worth N units, using as few coins as possible

DP approach

- table $c[1..n, 0..N]$ – one row for each denomination and one column for each amount from 0 units to N units
- $c[i,j]$ – minimum number of coins required to pay an amount j , with $0 \leq j \leq N$, using only coins of denominations 1 to i , $1 \leq i \leq n$
- the solution to the original problem is given by $c[n, N]$

Filling up the table

- $c[i,0] = 0$ for all i
- then the table can be filled row by row, left to right, or column by column, top to bottom
- to compute $c[i,j]$, two choices:
 - do not use any coins of denomination i , then $c[i,j] = c[i-1,j]$
 - use at least one coin of denomination i , then $c[i,j] = 1 + c[i,j - d_i]$
- therefore:
$$c[i,j] = \min(c[i-1,j], 1+c[i,j-d_i])$$
- if $i = 1$ and $j < d_1$, then set $c[i,j] = \infty$

DP algorithm

```
function coins(N)
  {Gives the minimum number of coins needed to make
   change for N units. Array d[1..n] specifies the coinage:
   in the example there are coins for 1, 4 and 6 units.}
  array d[1..n] = [1, 4, 6]
  array c[1..n, 0..N]
  for i ← 1 to n do c[i, 0] ← 0
  for i ← 1 to n do
    for j ← 1 to N do
      c[i, j] ← if i = 1 and j < d[i] then +∞
                  else if i = 1 then 1 + c[1, j - d[1]]
                  else if j < d[i] then c[i - 1, j]
                  else min(c[i - 1, j], 1 + c[i, j - d[i]])
  return c[n, N]
```

RT = $\Theta(nN)$

DP algorithm comments

- if an unlimited supply of coins of value 1 is available, then we can always find a solution to our problem
- if no coin with value 1, then there may be values of N for which no solution is possible
 - algorithm returns ∞

Finding how many coins of each denominations are used

PrintCoins(c,i,j)

```
if c[i,j] =  $\infty$ 
    then return no change possible
if j == 0, then return
if c[i,j] == c[i-1,j]
    then PrintCoins(c,i-1,j)
else if c[i,j] == 1 + c[i,j-di]
    then PrintCoins(c,i,j-di)
    print di
```

- Initial call: PrintCoins(c,n,N)
- RT = $O(N + n)$
- Example

0-1 knapsack problem

- Given:
 - n objects and a knapsack
 - for $i = 1, \dots, n$, object i has a positive weight w_i and a positive value v_i
 - objects may not be broken into smaller pieces, either take the whole object or leave it behind
 - the knapsack can carry a weight $\leq W$
- Objective: fill the knapsack s.t. to maximize the value of the included objects, while respecting the capacity constraints

DP approach

- table $V[1..n, 0..W]$ – one row for each available object and column for each weight from 0 to W
- $V[i, j]$ – maximum value of the objects we can transport if the weight limit is j , $0 \leq j \leq W$, and if we only include objects numbered from 1 to i , $1 \leq i \leq n$
- The solution to the original problem is given by $V[n, W]$

Filling up the table

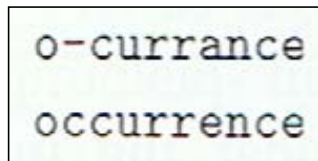
- $V[i,0] = 0$, for all i
- then the table can be filled row by row, left to right, or column by column, top to bottom
- to compute $V[i,j]$, two choices:
 - not adding object i to the knapsack, $V[i,j] = V[i-1, j]$
 - adding object i to the knapsack, $V[i,j] = v_i + V[i-1, j-w_i]$
- therefore:
$$V[i,j] = \max(V[i-1,j], v_i + V[i-1,j-w_i])$$
- for the out-of-bounds-entries, we define:
$$V[0,j] = 0 \text{ when } j \geq 0$$
$$V[i,j] = -\infty \text{ for all } i \text{ when } j < 0$$

Sequence Alignment – First example

- dictionaries on the Web and spell checkers
 - if you type “*ocurrance*”, it may ask: “Perhaps you mean *occurrence*?”
 - the dictionary will search its entries for the word most “similar” to the one typed
- how should we define *similarity* between two words or strings?

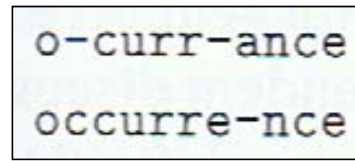
Modeling “similarity”

- model similarity between two strings by the number of *gaps* and *mismatches* that occur when lining up the two sequences



o-currance
occurrence

- one gap
- one mismatch



o-curr-ance
occurre-nce

- three gaps
- no mismatch

- which one is better?

Application – computational biology

- An organism's *genome*:
 - divided into linear DNA molecules, chromosomes, which serve as an one-dimensional chemical storage device
 - string over the alphabet {A,C,G,T} that determines the properties of the organism
 - adenine, cytosine, guanine, thymine
 - the string encodes instructions for building protein molecules
 - using a chemical mechanism to read portions of the chromosomes, a cell can construct proteins that control its metabolism

Application – computational biology

- let X and Y be two strains of bacteria, closely related evolutionary
- assume a certain substring in the DNA of X codes for a certain toxin
- if this substring is found in Y , we may hypothesize that this portion codes for a similar kind of toxin

Sequence Alignment Problem

- let $X = x_1x_2\ldots x_m$ and $Y = y_1y_2\ldots y_n$
- $\{1,2,\ldots,m\}$ and $\{1,2,\ldots,n\}$ represent different positions in strings X and Y
- *matching*: set of ordered pairs s.t. each item occurs in at most one pair
- a matching M of these two sets is an *alignment* if there are no *crossing* pairs: if $(i,j), (i',j') \in M$ and $i < i'$ then $j < j'$
 - an alignment provides a way to line up the two strings

stop- -tops

- alignment: $\{(2,1),(3,2),(4,3)\}$

Sequence Alignment Problem

Suppose M is a given alignment between X and Y

- *gap penalty*: each gap incurs a cost $\delta > 0$
- *mismatch cost*: for each pair of letters p, q in the alphabet, there is a mismatch cost α_{pq} for lining up p with q
 - usually $\alpha_{pp} = 0$
- the *cost* of M is the sum of gap and mismatch costs

Objective: find an ***optimal alignment***, that means an alignment of minimum cost.

- values δ and $\{\alpha_{pq}\}$ are given parameters
- the lower the cost, the more similar X and Y are
- going back to the first example *ocurrance* and *occurrence*, the first alignment is better if and only if $\delta + \alpha_{ae} < 3\delta$

Designing the DP algorithm

- in the optimal alignment M :
 - either $(m,n) \in M$ or $(m,n) \notin M$
- in any alignment M :
 - if $(m,n) \notin M$, then either the m^{th} position of X or the n^{th} position of Y are not matched in M
 - proof: assume by contradiction $(i,n), (m,j) \in M$ for some i,j .
Then $i < m$ and $j < n \Rightarrow$ crossing pairs \Rightarrow contradiction

Property: in any alignment M , one of the following is true:

- (i) $(m,n) \in M$, or
- (ii) the m^{th} position of X is not matched, or
- (iii) the n^{th} position of Y is not matched

Designing the DP algorithm

- let $\text{OPT}(i,j)$ – minimum cost of an alignment between $x_1x_2\dots x_i$ and $y_1y_2\dots y_j$

- recursively define $\text{OPT}(m,n)$:

case (i):

- pay $\alpha_{x_m y_n}$, then optimally align $x_1x_2\dots x_{m-1}$ and $y_1y_2\dots y_{n-1}$
 $\text{OPT}(m,n) = \alpha_{x_m y_n} + \text{OPT}(m-1,n-1)$

case (ii):

- pay gap cost δ , then optimally align $x_1x_2\dots x_{m-1}$ and $y_1y_2\dots y_n$
 $\text{OPT}(m,n) = \delta + \text{OPT}(m-1,n)$

case (iii):

- pay gap cost δ , then optimally align $x_1x_2\dots x_m$ and $y_1y_2\dots y_{n-1}$
 $\text{OPT}(m,n) = \delta + \text{OPT}(m,n-1)$

Designing the DP algorithm

Property: The minimum alignment costs satisfy the following recurrence, for $i \geq 1, j \geq 1$:

$$\text{OPT}(i,j) = \min\{\alpha_{x_i y_j} + \text{OPT}(i-1,j-1), \delta + \text{OPT}(i-1,j), \delta + \text{OPT}(i,j-1)\}$$

Moreover, (i,j) is in an optimal alignment iff the minimum is achieved by the first of these values.

Alignment (X,Y)

array $A[0..m,0..n]$

initialize array $A[i,0] = i\delta$ for each i

initialize array $A[0,j] = j\delta$ for each j

for $j = 1$ to n

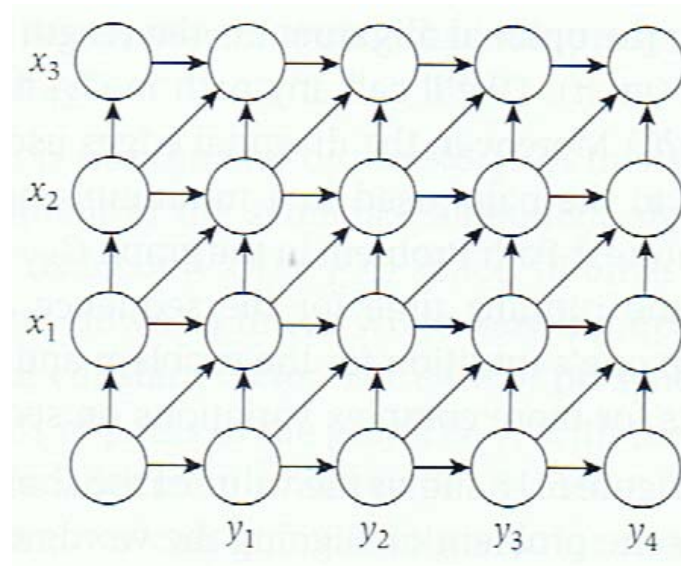
 for $i = 1$ to m

$A[i,j] = \min \{ \alpha_{x_i y_j} + A[i-1,j-1], \delta + A[i-1,j], \delta + A[i,j-1] \}$

return $A[m,n]$

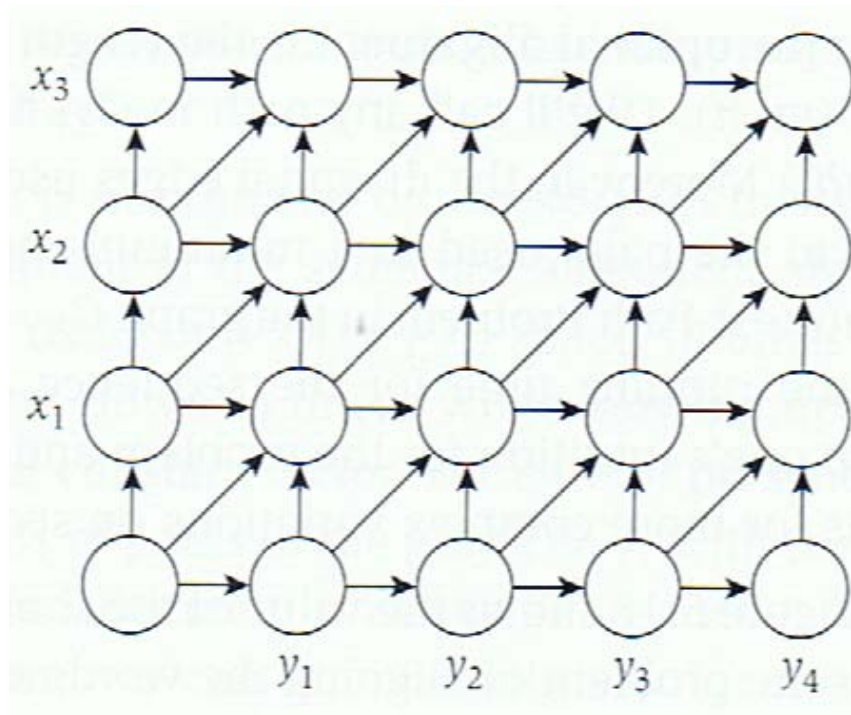
$$\text{RT} = \Theta(mn)$$

Graph – based picture of sequence alignment



- Cost of the edges:
 - horizontal & vertical edges have cost δ
 - diagonal edge $(i-1, j-1)$ to (i, j) has cost $\alpha_{x_i y_j}$
- Value of an optimal alignment is the minimum-cost of a path from $(0, 0)$ to (m, n)

Graph – based picture of sequence alignment



Let $f(i,j)$ denote the minimum cost of a path from $(0,0)$ to (i,j) in G_{xy} .
Then for all i,j , we have $f(i,j) = \text{OPT}(i,j)$.

Sequence Alignment in Linear Space via Divide and Conquer

- Finding the optimal alignment is equivalent to constructing the graph G_{XY} with mn nodes laid out in a grid and looking for the cheapest path between opposite corners
- RT for this approach is $O(mn)$ and space $O(mn)$
- This is too large for biological applications where strings are very long
 - if the two strings have $\sim 100,000$ symbols each, then RT ~ 10 billion primitive operations and space is ~ 10 billion array
- Objective: enhancement of the sequence alignment algorithm that has RT = $O(mn)$ and space $O(m+n)$
 - uses divide-and-conquer

Designing the Algorithm

- First, we'll show that if we care only about the *value* of an optimal alignment, then it's easy to have linear space
- Key observation: to fill out the array A, we only need information on the current column of A and the previous column of A
- Instead of using array A of size $(m+1) \times (n+1)$, use array B of size $(m+1) \times 2$
- As the algorithm iterates through values of j, entries $B[i,0]$ will hold the *previous* column's value $A[i,j-1]$ and entries of the form $B[i,1]$ will hold the *current* column's values $A[i,j]$

Designing the Algorithm

Space-Efficient-Alignment (X,Y)

```
array B[0..m,0..1]
initialize B[i,0] =  $i\delta$  for each i (just as in column 0 of A)
for j = 1,...,n
    B[0,j] =  $j\delta$  (since this corresponds to entry A[0,j])
    for i = 1,...,m
        B[i,1] = min[  $\alpha_{x_i y_j} + B[i-1,0]$ ,  $\delta + B[i-1,1]$ ,  $\delta + B[i,0]$  ]
    endfor
    // move col 1 of B to col 0 to make room for the next iteration
    update B[i,0] = B[i,1] for each i
endfor
```

- $RT = O(mn)$
- $space = O(m)$

Space-efficient-alignment

- when the alg. terminates, $B[i,1]$ holds the value of $OPT(i,n)$ for $i = 0, \dots, m$
- issue: how to find the assignment itself?
 - we haven't left enough information to find the alignment
 - B has only the last two columns, so we cannot trace back the optimal alignment (shortest path)
- different approach if we want to recover the optimal alignment

A backward formulation of the DP

- $f(i,j)$ – length of the shortest path $(0,0)$ to (i,j) in the graph G_{XY}
 $f(i,j) = \text{OPT}(i,j)$
- define $g(i,j)$ – length of the shortest path from (i,j) to (m,n) in G_{XY}
- build g using DP in reverse: start with $g(m,n) = 0$, and the answer we want is $g(0,0)$

- for $i < m$ and $j < n$ we have:

$$g(i,j) = \min [\alpha_{x_{i+1} y_{j+1}} + g(i+1, j+1), \delta + g(i, j+1), \delta + g(i+1, j)]$$

- g is built using DP backward from (m,n)
- we can also design the space-efficient version,
Backward-Space-Efficient-Alignment(X, Y)
in space $O(m)$ and $\text{RT} = O(mn)$

Combining the Forward and Backward Formulations

Important properties:

- The length of the shortest corner-to-corner path in G_{XY} that passes through (i,j) is $f(i,j) + g(i,j)$
- Let k be any number in $\{0, \dots, n\}$ and let q be an index minimizes the quantity $f(q,k) + g(q,k)$. Then there is a corner-to-corner path of minimum length that passes through the node (q,k) .

Designing the algorithm

- divide G_{XY} along the center column and compute $f(i, n/2)$ and $g(i, n/2)$ for each i , using the two space-efficient algorithms
- find the minimum $f(i, n/2) + g(i, n/2)$ for same value i
- then there is a shortest corner-to-corner path that passes through $(i, n/2)$
- recursively find the shortest-path in G_{XY} between $(0,0)$ and $(i, n/2)$ and in the portion between $(i, n/2)$ and (m,n)
- MAIN IDEA:
 - Apply these recursive calls sequentially and reuse the working space from one call to the next
- then the space usage is $O(m+n)$

Designing the algorithm

- maintain a globally accessible list P with nodes on the shortest corner-to-corner path as they are discovered
 - initially, P is empty
 - P has at most $m+n$ entries, since any path has at most $m+n$ edges
- notation:
 $X[i:j]$, for $1 \leq i \leq j \leq m$, is the substring $x_i x_{i+1} \dots x_j$
similar for $Y[i:j]$
- assume for simplicity n is a power of 2

Designing the algorithm

Divide-and-Conquer-Alignment(X,Y)

m is the number of symbols in X

n is the number of symbols in Y

if $m \leq 2$ or $n \leq 2$ then

 compute optimal alignment using Alignment(X,Y)

call Space-Efficient-Alignment(X,Y[1:n/2])

call Backward-Space-Efficient-Alignment(X,Y[n/2+1:n])

let q be the index minimizing $f(q,n/2) + g(q,n/2)$

add (q,n/2) to the global list P

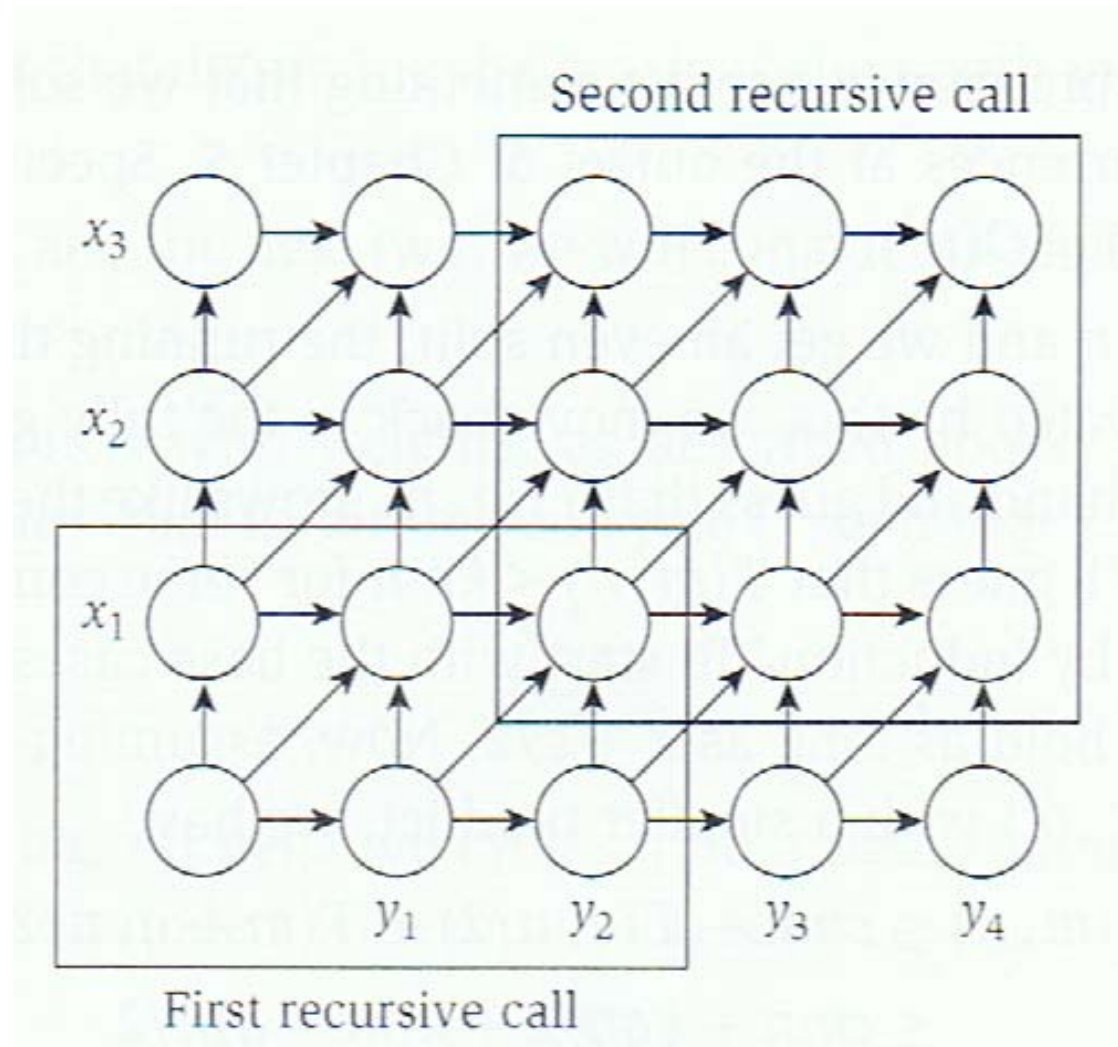
Divide-and-Conquer-Alignment(X[1:q],Y[1:n/2])

Divide-and-Conquer-Alignment(X[q+1:m],Y[n/2+1:n])

return P

✓ Space used is $O(m+n)$

Example



RT analysis

The RT of Divide-and-Conquer-Alignment on strings of length m and n is $O(mn)$.

Proof:

$T(m,n)$ – running time

$$T(m,n) \leq cmn + T(q,n/2) + T(m-q,n/2)$$

$$T(m,2) \leq cm$$

$$T(2,n) \leq cn$$

Particular case: $m = n$ and q is in the middle

$$T(n) \leq cn^2 + 2T(n/2)$$

case 3 of the Master Thm $\Rightarrow T(n) = \Theta(n^2)$

RT analysis

General case:

$$T(m,n) \leq cmn + T(q,n/2) + T(m-q,n/2)$$

Show by induction that $T(m,n) = O(mn)$, that means

$$T(m,n) \leq kmn \text{ for some constant } k$$

Base case: $m \leq 2$ or $n \leq 2$ is true

Inductive step:

$$\begin{aligned} T(m,n) &\leq cmn + T(q,n/2) + T(m-q,n/2) \\ &\leq cmn + kqn/2 + k(m-q)n/2 \\ &= cmn + kqn/2 + kmn/2 - kqn/2 \\ &= (c + k/2)mn \end{aligned}$$

Inductive step works if $c+k/2 = k$, that means $c = k/2$ or $k = 2c$